

# EVALUATING MODEL FIT

---

## EVALUATING MODEL FIT

---

# LEARNING OBJECTIVES

- Define regularization, bias, and error metrics for regression problems
- Evaluate model fit using loss functions
- Select regression methods based on fit and complexity

---

**COURSE**

---

**PRE-WORK**

---

# PRE-WORK REVIEW

---

- Understand goodness of fit (r-squared)
- Measure statistical significance of features
- Recall what a residual is
- Implement a sklearn estimator to predict a target variable

---

**OPENING**

---

# R-SQUARES AND RESIDUALS

---

## WHAT IS R-SQUARED? WHAT IS A RESIDUAL?

---

- R-squared, the central metric introduced for linear regression
- Which model performed better, one with an r-squared of 0.79 or 0.81?
- R-squared measures explain variance.
- But does it tell the magnitude or scale of error?
- We'll explore loss functions and find ways to refine our model.


## **INTRODUCTION**

---

# **LINEAR MODELS AND ERROR**

## RECALL: WHAT'S RESIDUAL ERROR?

- ▶ In linear models, residual error must be normal with a median close to zero.
- ▶ Individual residuals are useful to see the error of specific points, but it doesn't provide an overall picture for optimization.
- ▶ We need a metric to summarize the error in our model into one value.
- ▶ Mean square error: the mean residual error in our model





---

## MEAN SQUARED ERROR (MSE)

---

- To calculate MSE:
  - Calculate the difference between each target  $y$  and the model's predicted value  $\hat{y}$  (i.e. the residual)
  - Square each residual.
  - Take the mean of the squared residual errors.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$

---

## MEAN SQUARED ERROR (MSE)

---

- sklearn's metrics module includes a mean\_squared\_error function.

```
from sklearn import metrics
```

```
metrics.mean_squared_error(y, model.predict(X))
```

---

## MEAN SQUARED ERROR (MSE)

---

- For example, two arrays of the same values would have an MSE of 0.

```
from sklearn import metrics
```

```
metrics.mean_squared_error([1, 2, 3, 4, 5], [1, 2, 3, 4, 5])
```

```
0.0
```

---

## MEAN SQUARED ERROR (MSE)

---

- Two arrays with different values would have a positive MSE.

```
from sklearn import metrics
metrics.mean_squared_error([1, 2, 3, 4, 5], [5, 4, 3, 2, 1])
# (4^2 + 2^2 + 0^2 + 2^2 + 4^2) / 5
8.0
```

---

## HOW DO WE MINIMIZE ERROR?

---

- The regression method we've used is called “Ordinary Least Squares”.
- This means that given a matrix  $X$ , solve for the *least* amount of square error for  $y$ .
- However, this assumes that  $X$  is unbiased, that it is representative of the population.

---

# LET'S COMPARE TWO RANDOM MODELS

---

```
import numpy as np
import pandas as pd
from sklearn import linear_model
```

```
df = pd.DataFrame({'x': range(100), 'y': range(100)})
biased_df = df.copy()
biased_df.loc[:20, 'x'] = 1
biased_df.loc[:20, 'y'] = 1
```

```
def append_jitter(series):
    jitter = np.random.random_sample(size=100)
    return series + jitter
```

---

# LET'S COMPARE TWO RANDOM MODELS

---

```
df['x'] = append_jitter(df.x)
df['y'] = append_jitter(df.y)
```

```
biased_df['x'] = append_jitter(biased_df.x)
biased_df['y'] = append_jitter(biased_df.y)
```

- Fit:


```
lm = linear_model.LinearRegression().fit(df[['x']], df['y'])
print metrics.mean_squared_error(df['y'], lm.predict(df[['x']]))
```

- Biased fit:


```
lm = linear_model.LinearRegression().fit(biased_df[['x']], biased_df['y'])
print metrics.mean_squared_error(df['y'], lm.predict(df[['x']]))
```

# BIAS VS. VARIANCE


---




High bias  
High variance



Low bias  
High variance



High bias  
Low variance



Low bias  
Low variance




---

# BIAS VARIANCE TRADEOFF

---

- When our error is ***biased***, it means the model's prediction is consistently far away from the actual value.
- This could be a sign of poor sampling and poor data.
- One objective of a biased model is to trade bias error for generalized error. We prefer the error to be more evenly distributed across the model.
- This is called error due to ***variance***.
- We want our model to *generalize* to data it hasn't seen even if doesn't perform as well on data it has already seen.

# BIAS VARIANCE TRADEOFF



---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS (5 minutes)

1. Which of the following scenarios would be better for a weatherman?
  - a. Knowing that I can very accurately "predict" the temperature outside from previous days perfectly, but be 20-30 degrees off for future days
  - b. Knowing that I can accurately predict the general trend of the temperate outside from previous days, and therefore am at most only 10 degrees off on future days

### DELIVERABLE

Answers to the above questions

**DEMO**

---

# CROSS VALIDATION

---

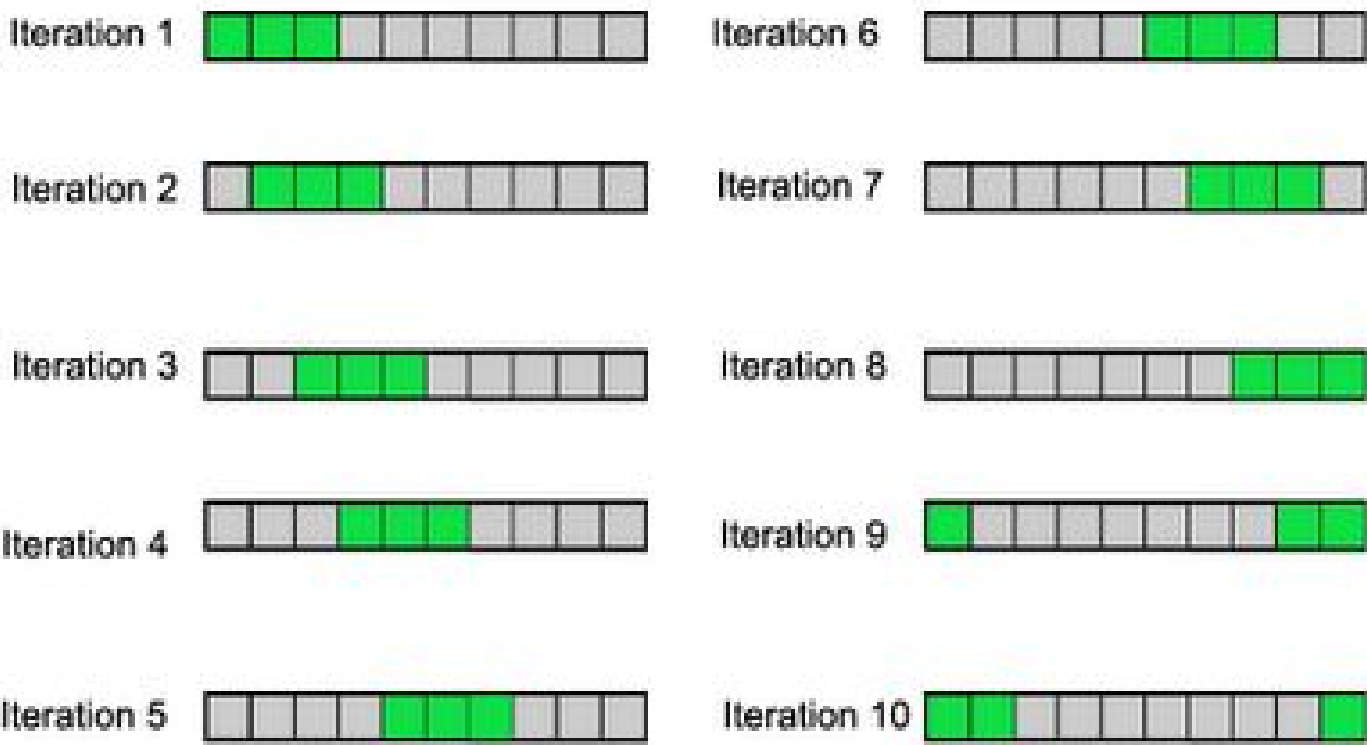
# CROSS VALIDATION

---

- Cross validation can help account for bias.
- The general idea is to
  - Generate several models on different cross sections of the data
  - Measure the performance of each
  - Take the mean performance
- This technique swaps bias error for generalized error, describing previous trends accurately enough to extend to future trends.

# CROSS VALIDATION

INPUT PARAMETERS:  
Number of iterations = ~~11~~ 10       $\Delta = 1$   
% Train = 30%      train Instances = 3



Train Instances

Test Instances

---

# K-FOLD CROSS VALIDATION

---

- k-fold cross validation
  - Split the data into  $k$  group
  - Train the model on all segments except one
  - Test model performance on the remaining set
- If  $k = 5$ , split the data into five segments and generate five models.

---

# USING K-FOLD CROSS VALIDATION WITH MSE

---

- Import the appropriate packages and load data.

```
from sklearn import cross_validation
wd = '../..../datasets/'
bikeshare = pd.read_csv(wd + 'bikeshare/bikeshare.csv')
weather = pd.get_dummies(bikeshare.weathersit, prefix='weather')
modeldata = bikeshare[['temp', 'hum']].join(weather[['weather_1',
'weather_2', 'weather_3']])
y = bikeshare.casual
```



---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

- Build models on subsets of the data and calculate the average score.

```
kf = cross_validation.KFold(len(modeldata), n_folds=5, shuffle=True)
scores = []
for train_index, test_index in kf:
    lm =
linear_model.LinearRegression().fit(modeldata.iloc[train_index],
y.iloc[train_index])
    scores.append(metrics.mean_squared_error(y.iloc[test_index],
lm.predict(modeldata.iloc[test_index])))

print np.mean(scores)
```

---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

▸ This can be compared to the model built on all of the data.

- This score will be lower, but we're trading off bias error for generalized error:

```
lm = linear_model.LinearRegression().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
```

▸ Which approach would predict new data more accurately?

## **GUIDED PRACTICE**

---

# **CROSS VALIDATION WITH LINEAR REGRESSION**

---

# ACTIVITY: CROSS VALIDATION WITH LINEAR REGRESSION

---



## EXERCISE

### DIRECTIONS (20 minutes)

If we were to continue increasing the number of folds in cross validation, would error increase or decrease?

1. Using the previous code example, perform k-fold cross validation for all even numbers between 2 and 50.
2. Answer the following questions:
  - a. What does `shuffle=True` do?
  - b. At what point does cross validation no longer seem to help the model?
3. Hint: `range(2, 51, 2)` produces a list of even numbers from 2 to 50

### DELIVERABLE

Answers to questions

## **INTRODUCTION**

---

# **REGULARIZATION AND CROSS VALIDATION**

---

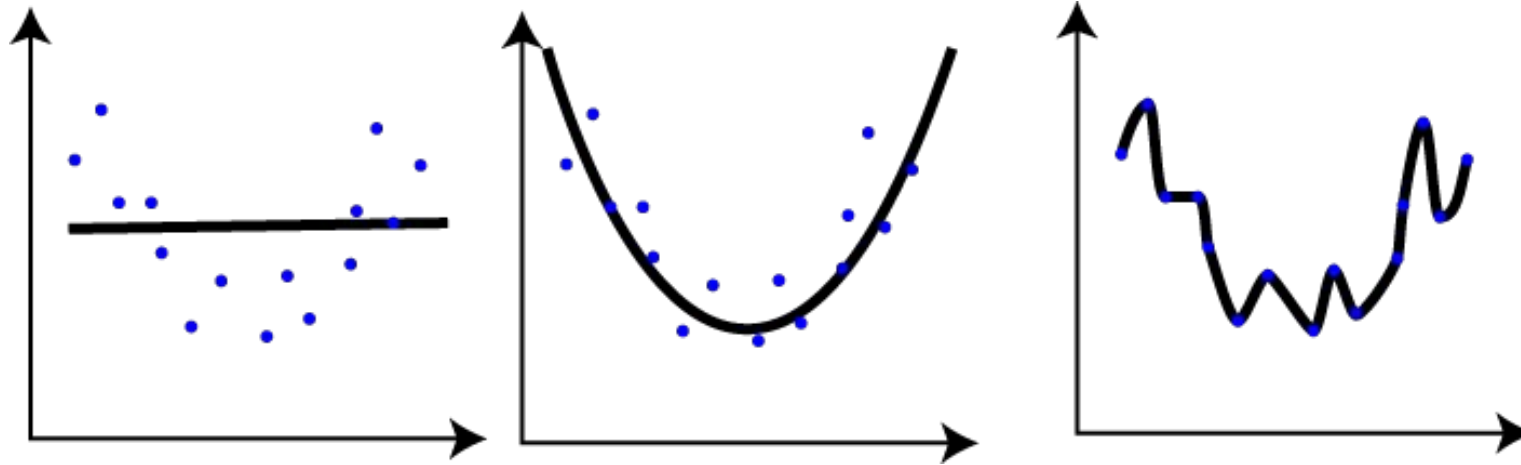
## WHAT IS REGULARIZATION? AND WHY DO WE USE IT?

---

- Regularization is an additive approach to protect models against overfitting (being potentially biased and overconfident, not generalizing well).
- Regularization becomes an additional weight to coefficients, shrinking them closer to zero.
- L1 (Lasso Regression) adds the extra weight to coefficients.
- L2 (Ridge Regression) adds the square of the extra weight to coefficients.
- Use Lasso when we have more features than observations ( $k > n$ ) and Ridge otherwise.

# WHAT IS OVERFITTING?

---



- The first model poorly explains the data.
- The second model explains the general curve of the data.
- The third model drastically overfits the model, bending to every point.
- Regularization helps prevent the third model.

---

## WHERE REGULARIZATION MAKES SENSE

---

- What happens to MSE if use Lasso or Ridge Regression directly?

```
lm = linear_model.LinearRegression().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
lm = linear_model.Lasso().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
lm = linear_model.Ridge().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
```

```
1672.58110765 # OLS
1725.41581608 # L1
1672.60490113 # L2
```



---

## WHERE REGULARIZATION MAKES SENSE

---

- It doesn't seem to help. Why is that?
- We need to optimize the regularization weight parameter (called  $\alpha$ ) through cross validation.

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS (5 minutes)

1. Why is regularization important?
2. What does it protect against and how?

### DELIVERABLE

Answers to the above questions

**DEMO**

---

# UNDERSTANDING REGULARIZATION EFFECTS

---

## QUICK CHECK

---

- We are working with the bikeshare data to predict riders over hours/days with a few features.
- Does it make sense to use a ridge regression or a lasso regression?
- Why?

---

# UNDERSTANDING REGULARIZATION EFFECTS

---

- Let's test a variety of alpha weights for Ridge Regression on the bikeshare data.

```
alphas = np.logspace(-10, 10, 21)
for a in alphas:
    print 'Alpha:', a
    lm = linear_model.Ridge(alpha=a)
    lm.fit(modeldata, y)
    print lm.coef_
    print metrics.mean_squared_error(y, lm.predict(modeldata))
```

- What happens to the weights of the coefficients as alpha increases? What happens to the error as alpha increases?

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

- Grid search exhaustively searches through all given options to find the best solution. Grid search will try all combos given in `param_grid`.

```
param_grid = {  
    'intercept': [True, False],  
    'alpha': [1, 2, 3],  
}
```

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

▸ This param grid has six different options:

- intercept True, alpha 1
- intercept True, alpha 2
- intercept True, alpha 3
- intercept False, alpha 1
- intercept False, alpha 2
- intercept False, alpha 3

```
param_grid = {  
    'intercept': [True, False],  
    'alpha': [1, 2, 3],  
}
```

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

- This is an incredibly powerful, automated machine learning tool!

```
from sklearn import grid_search
```

```
alphas = np.logspace(-10, 10, 21)  
gs = grid_search.GridSearchCV(  
    estimator=linear_model.Ridge(),  
    param_grid={'alpha': alphas},  
    scoring='mean_squared_error')
```



---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

```
gs.fit(modeldata, y)
```

```
print -gs.best_score_ # mean squared error here comes in negative, so  
let's make it positive.
```

```
print gs.best_estimator_ # explains which grid_search setup worked  
best
```

```
print gs.grid_scores_ # shows all the grid pairings and their  
performances.
```

## GUIDED PRACTICE

---

# GRID SEARCH CV, SOLVING FOR ALPHA

---

# ACTIVITY: GRID SEARCH CV, SOLVING FOR ALPHA

---



## EXERCISE

### DIRECTIONS (25 minutes)

1. Modify the previous code to do the following:
  - a. Introduce cross validation into the grid search. This is accessible from the cv argument.
  - b. Add `fit_intercept = True` and `False` to the `param_grid` dictionary.
  - c. Re-investigate the best score, best estimator, and grid score attributes as a result of the grid search.

### DELIVERABLE

New code and output that meets above requirements

## INTRODUCTION

---

# MINIMIZING LOSS THROUGH GRADIENT DESCENT


---

# GRADIENT DESCENT

---

- Gradient Descent can also help us minimize error.
- How Gradient Descent works:
  - A random linear solution is provided as a starting point
  - The solver attempts to find a next “step”: take a step in any direction and measure the performance.
  - If the solver finds a better solution (i.e. lower MSE), this is the new starting point.
  - Repeat these steps until the performance is optimized and no “next steps” perform better. The size of steps will shrink over time.

# GRADIENT DESCENT



---

# A CODE EXAMPLE OF GRADIENT DESCENT

---

```
num_to_approach, start, steps, optimized = 6.2, 0., [-1, 1], False
while not optimized:
    current_distance = num_to_approach - start
    got_better = False
    next_steps = [start + i for i in steps]
    for n in next_steps:
        distance = np.abs(num_to_approach - n)
        if distance < current_distance:
            got_better = True
            print distance, 'is better than', current_distance
            current_distance = distance
            start = n
```

---

# A CODE EXAMPLE OF GRADIENT DESCENT

---

```
if got_better:  
    print 'found better solution! using', current_distance  
    a += 1  
else:  
    optimized = True  
    print start, 'is closest to', num_to_approach
```

- What is the code doing? What could go wrong?



# GLOBAL VS LOCAL MINIMUMS

---

- Gradient Descent could solve for a *local* minimum instead of a *global* minimum.
- A *local* minimum is confined to a very specific subset of solutions. The *global* minimum considers all solutions. These could be equal, but that's not always true.



**DEMO**

---

# APPLICATION OF GRADIENT DESCENT

---

# APPLICATION OF GRADIENT DESCENT

---

- Gradient Descent works best when:
  - We are working with a large dataset. Smaller datasets are more prone to error.
  - Data is cleaned up and normalized.
- Gradient Descent is significantly faster than OLS. This becomes important as data gets bigger.

---

# APPLICATION OF GRADIENT DESCENT

---

- We can easily run a Gradient Descent regression.
- Note: The verbose argument can be set to 1 to see the optimization steps.

```
lm = linear_model.SGDRegressor()
```

```
lm.fit(modeldata, y)
```

```
print lm.score(modeldata, y)
```

```
print metrics.mean_squared_error(y, lm.predict(modeldata))
```

- Untuned, how well did gradient descent perform compared to OLS?

---

# APPLICATION OF GRADIENT DESCENT

---

- Gradient Descent can be tuned with
  - the learning rate: how aggressively we solve the problem
  - epsilon: at what point do we say the error margin is acceptable
  - iterations: when should we stop no matter what

---

**INDEPENDENT PRACTICE**

---

**ON YOUR OWN**

---

# ACTIVITY: ON YOUR OWN

---



## EXERCISE

### DIRECTIONS (30 minutes)

There are tons of ways to approach a regression problem.

1. Implement the Gradient Descent approach to our bikeshare modeling problem.
2. Show how Gradient Descent solves and optimizes the solution.
3. Demonstrate the `grid_search` module.
4. Use a model you evaluated last class or the simpler one from today. Implement `param_grid` in grid search to answer the following questions:
  - a. With a set of values between  $10^{-10}$  and  $10^{-1}$ , how does MSE change?
  - b. Our data suggests we use L1 regularization. Using a grid search with `l1_ratios` between 0 and 1, increasing every 0.05, does this statement hold true? If not, did gradient descent have enough iterations to work properly?
  - c. How do these results change when you alter the learning rate?

### DELIVERABLE

Gradient Descent approach and answered questions

# ACTIVITY: ON YOUR OWN



## EXERCISE

### Starter Code

```
params = {} # put your gradient descent parameters here
gs = grid_search.GridSearchCV(
    estimator=linear_model.SGDRegressor(),
    cv=cross_validation.KFold(len(modeldata), n_folds=5, shuffle=True),
    param_grid=params,
    scoring='mean_squared_error',
)

gs.fit(modeldata, y)

print 'BEST ESTIMATOR'
print -gs.best_score_
print gs.best_estimator_
print 'ALL ESTIMATORS'
print gs.grid_scores_
```



---

**CONCLUSION**

---

# TOPIC REVIEW

---

## LESSON REVIEW

---

- What's the (typical) range of  $r$ -squared?
- What's the range of mean squared error?
- How would changing the scale or interpretation of  $y$  (your target variable) effect mean squared error?
- What's cross validation, and why do we use it in machine learning?
- What is error due to bias? What is error due to variance? Which is better for a model to have, if it had to have one?
- How does gradient descent try a different approach to minimizing error?

**COURSE**

---

**BEFORE NEXT CLASS**

---

**LESSON**

---

**CREDITS**

---

**LESSON**

---

**Q & A**

## **LESSON**

---

# **EXIT TICKET**

**DON'T FORGET TO FILL OUT YOUR EXIT TICKET**